

Should Agents Reason Linearly or as a Tree?

Baseline (IO) vs. Self-Refine vs. Tree-of-Thought vs. Agentic-Program-Repair

Same models, same prompts, same tasks. Different reasoning loop.

Methods

Method	Paper	What it does
Baseline (IO)	—	One-shot autoregressive rollout
Self-Refine	Madaan et al. 2023	LLM critiques its own answer at every round
Tree-of-Thought	Yao et al. 2023	Keeps several half-finished attempts alive, scores them, expands the best, prunes the rest.
Agentic-Program-Repair	Maddila et al. 2025	Make N independent attempts, keep those that pass tests.

Summary: Agentic Program Repair

Use deterministic, verifiable tests to determine correctness. Think: Test-Driven Development.

In the paper: an agent reads the failing tests, edits code, runs the tests, and repeats. Their headline result is not a clever algorithm — it is **feedback plus retries**:

Their system	Solve rate
Agent alone	28.5%
+ static analysis feedback	34.1%
+ test execution feedback	43.9%
Best single patch of many attempts	46.3%
5 repeated runs of the whole loop	61.0%

The Engineering Agent is an internal Meta system and is **not open-source**, so our implementation is built from the paper alone — keeping the transferable part: fresh attempts, judged by tests.

Our Experiments

Family	Task	The model must produce	How we grade it
Game of 24	Use four numbers and + - * / to reach 24	Equations + an Answer: line	Exact arithmetic check
Coding	Write a function from a docstring	A code block	Real unit tests in a subprocess
Letter counting	Count one letter in a tricky word	Inventory line + count	Exact count


Nothing is graded by a model. Every pass/fail comes from arithmetic, a test run, or a string compare.

Experiment Setup

Models	llama2:7b-chat and mistral:7b-instruct — two 7B models from the same era as the papers
Serving	Ollama, 2 GPUs, NUM_PARALLEL=4 , 6 tasks in flight
Scale	2 models × 4 methods × 17 tasks = 136 runs
Budget	Same prompts for all four methods; 60 model calls per task maximum
Metrics	Solve rate · generated tokens · model calls · wall-clock time


Cost is reported per *solved* task — total tokens spent divided by tasks actually solved.

Success rate




Findings: Agentic-PR leads overall (35%) and on both coding and counting. Tree-of-Thought only works on the stronger model (53% Mistral vs 6% Llama-2). Game of 24 is a wall: 1 solve in 48.

Tokens spent



Findings: Tree-of-Thought costs 9× more per solve on Llama-2 than on Mistral (14,045 vs 1,524), same code. Agentic-PR is the cheapest route to accuracy.

Wall-clock time



Findings: Baseline 2.1s per run, Agentic-PR 6.1s, Self-Refine 10.0s, Tree-of-Thought 11.2s. Time tracks how many calls a method makes — not how good it is.

What I take away

1. **Verification is the win, not the tree.** Tests turn extra compute into correct answers. More thinking without a check mostly buys longer wrong answers.
2. **The methods are not interchangeable.** Agentic-PR owns counting and coding; Tree-of-Thought's best case is a harder task on a stronger model.
3. **Self-Refine is the weakest link here.** Most expensive, least accurate — worse than one-shot on every family.
4. **Time tracks the loop, not the result.** Baseline 2.1s per run, Agentic-PR 6.1s, Self-Refine 10.0s, Tree-of-Thought 11.2s — and the extra time does not buy accuracy for the two model-driven methods.
5. **Cost depends on the model as much as the method.** Tree-of-Thought cost 9× more per solve on Llama-2 than on Mistral, with identical code.

Caveats: 34 runs per method is directional, not statistical. Two 7B models from 2023. Game of 24 was out of reach, so the math family contributes almost nothing. Timings include GPU queue time.

Appendix A — Task distribution

Family	Tasks	Task IDs
Game of 24	6	g24-0 ... g24-5
Coding	5	prog-merge_intervals , prog-rotate , prog-kth_largest , prog-parse_formula , prog-merge_sorted_arrays
Letter counting	6	count-strawberry , count-raspberry , count-Mississippi , count-bookkeeper , count-dreadnought , count-refrigerator
Total	17	

Each task gets 8 runs (2 models × 4 methods) → $2 \times 4 \times 17 = 136$ runs.

Six tasks were solved by nobody: g24-1 ... g24-5 and prog-parse_formula .

Appendix B — What APR is in our implementation

Independent fresh attempts, verified by tests. Not self-refine, and not a combination.

Attempt	What the model sees
1	the problem prompt, alone
2	the problem prompt + the rejected answer + <i>"start again with a DIFFERENT approach"</i>
3	the same, with attempt 2 as the rejected answer

Each attempt is a complete, independent answer. Test results decide **whether to retry** — nothing is carried forward as a partial solution.

Honest gap vs. the paper: we tell the model *"the verifier rejected it"*, not *which test failed*. APR feeds the real test output back into the loop, so our signal is weaker than theirs.

Appendix C — Key parameters (1 of 2)

Parameter	Value
Models	llama2:7b-chat , mistral:7b-instruct (7B, 2023-era)
Serving	Ollama 0.40.0, CUDA, 2 × RTX PRO 6000 (95 GB each)
Concurrency	NUM_PARALLEL=4 , 6 runs in flight (--jobs 6)
Tasks	17 per model per method (6 math / 5 coding / 6 counting)
Runs	2 models × 4 methods × 17 tasks = 136
Call cap	60 model calls per task (max observed: 10)
Output cap	450 tokens per generation

Appendix C — Key parameters (2 of 2)

Method	Setting
Baseline (IO)	temperature 0.0, 1 call
Self-Refine	temperature 0.7, max 4 rounds; critique at temperature 0.0, 120 tokens
Agentic-Program-Repair	temperature 0.9, 3 independent attempts
Tree-of-Thought	temperature 0.8; depth 3 (math) / 2 (others); expand 2; beam 2; scoring at temperature 0.0

Verifiers (`tasks.py`): unit tests in a subprocess with a 10 s timeout; an exhaustive solver for Game of 24; exact matching for counts.

Deliberate fairness choice: `--iters 4` (Self-Refine's paper default) and `--candidates 3` give all methods comparable numbers of attempts.